

## ECEn 528

### Study Guide - Dynamic Scheduling and Speculation

- Read Sections C.7, 3.4-3.6 and 3.8 of H&P
  - Things to focus on
    - Do not worry too much about all of the details of how the scoreboard maintains instruction state in section C.7; instead, focus on the basic idea that the scoreboard must look for potential WAR, WAW, and RAW hazards and prevent instructions from issuing when the hazards are there
    - Again, don't worry about all the details of Tomasulo's scheme in Figure 3.9. But do understand why it works: how the WAR/WAW hazards are removed and the essential dependences satisfied. The loop-based example is particularly important.
    - Working through Case Study 1 in chapter 3 as we learn the techniques could be very helpful
    - Don't worry about all the details in figure 3.14, but do understand how the reorder buffer actually works
    - Understand why there is a difference in timing between figures 3.19 and 3.20
    - If you're interested in history, here is Robert Tomasulo talking about the IBM 360/91 at University of Michigan: <http://leccap.engin.umich.edu/leccap/view/tomasulo/5047>
  - Clarifications
    - Dynamic scheduling is also known as out-of-order execution (OoO)
    - Dynamic scheduling is an interesting variation of “clever” solutions to hazards: we still “stall” for hazards, but the stall of one instruction doesn't affect other instructions directly as it does in an in-order pipe.
    - The description of Tomasulo's scheme uses dynamic scheduling only for the floating-point units; the algorithm can be used for all instructions. Historically, though, when it was first used, it was only used for floating-point instructions.
    - The book says that placing the instruction in a reservation station is “issue”, and then starting its execution is “execute”. The alternate terminology they suggest (“dispatch” and “issue”) is also extremely common, and more descriptive. Even worse, some people mix the two and call the process “issue” then “dispatch”!
    - The description of exception handling on p. 174 does **not** provide precise exceptions!
    - The description of reservation stations assumes one “group” of reservation stations per functional unit; it is possible for functional units to share a reservation station, though the control logic is more complex.
    - A statement is made on page 175 that in a dynamically scheduled pipeline, there must be one extra cycle of latency between a producing instruction and a consuming instruction. This is not true in general, but does require more hardware to remove that cycle and the circuit timing is tricky. However, as we will see in a paper we read later on, removing that cycle is very important.
    - The subtle distinction made in the second paragraph of Section 3.6 is that their presentation of dynamic scheduling did not allow scheduling beyond a branch and didn't provide precise exceptions
    - Note that a “dataflow execution” where instructions execute whenever their operands are available quite naturally leads to the issue of multiple instructions per cycle.
    - The reorder buffer is like a future file, but is not organized like a register file, but rather a FIFO with associative lookup (on register numbers)
  - Answer the following questions:
    1. Why are WAR hazards possible in a dynamically scheduled pipeline?

2. Why is dynamic scheduling helpful? Think of a situation in which dynamic scheduling would perform better than static scheduling.
3. Why do compilers still schedule code even if the hardware uses dynamic scheduling?
4. Why does Tomasulo's scheme allow removal of hazards which arise from name dependences which a compiler cannot?
5. Why does the description of exception handling on p. 95-96 not provide precise exceptions?
6. Why does dynamic scheduling require accurate branch prediction?
7. How does the use of a reorder buffer differ from Tomasulo's scheme?
8. Why must commit be performed in order?

9. How many instructions can be in flight when using a reorder buffer? How many when using register renaming?
10. Why might you want to handle a TLB miss at commit instead of when it is first detected?
11. How would exceptions or branch mispredictions be handled in a register renaming scheme?